

Docker, Clustering, Scalability

A Case Study

John Quick

Smartvox Limited

The Design Specification



SIP Registrar/Proxy in front of a
FreeSwitch-based core



Docker-based (Bare metal hosting)



Scalable (Microservices concept)



Short Term Plans: WebRTC support
Deploy using Kubernetes



Longer Term Plans: Multi site deployment

Strategy



Identify gaps in my knowledge



Search the web for ideas - reports, articles, blogs



Read product documentation



Set up a test lab



Learn about concepts through testing



Identify issues and snags through testing

Key Topics

Docker


Clustering topologies

Load sharing

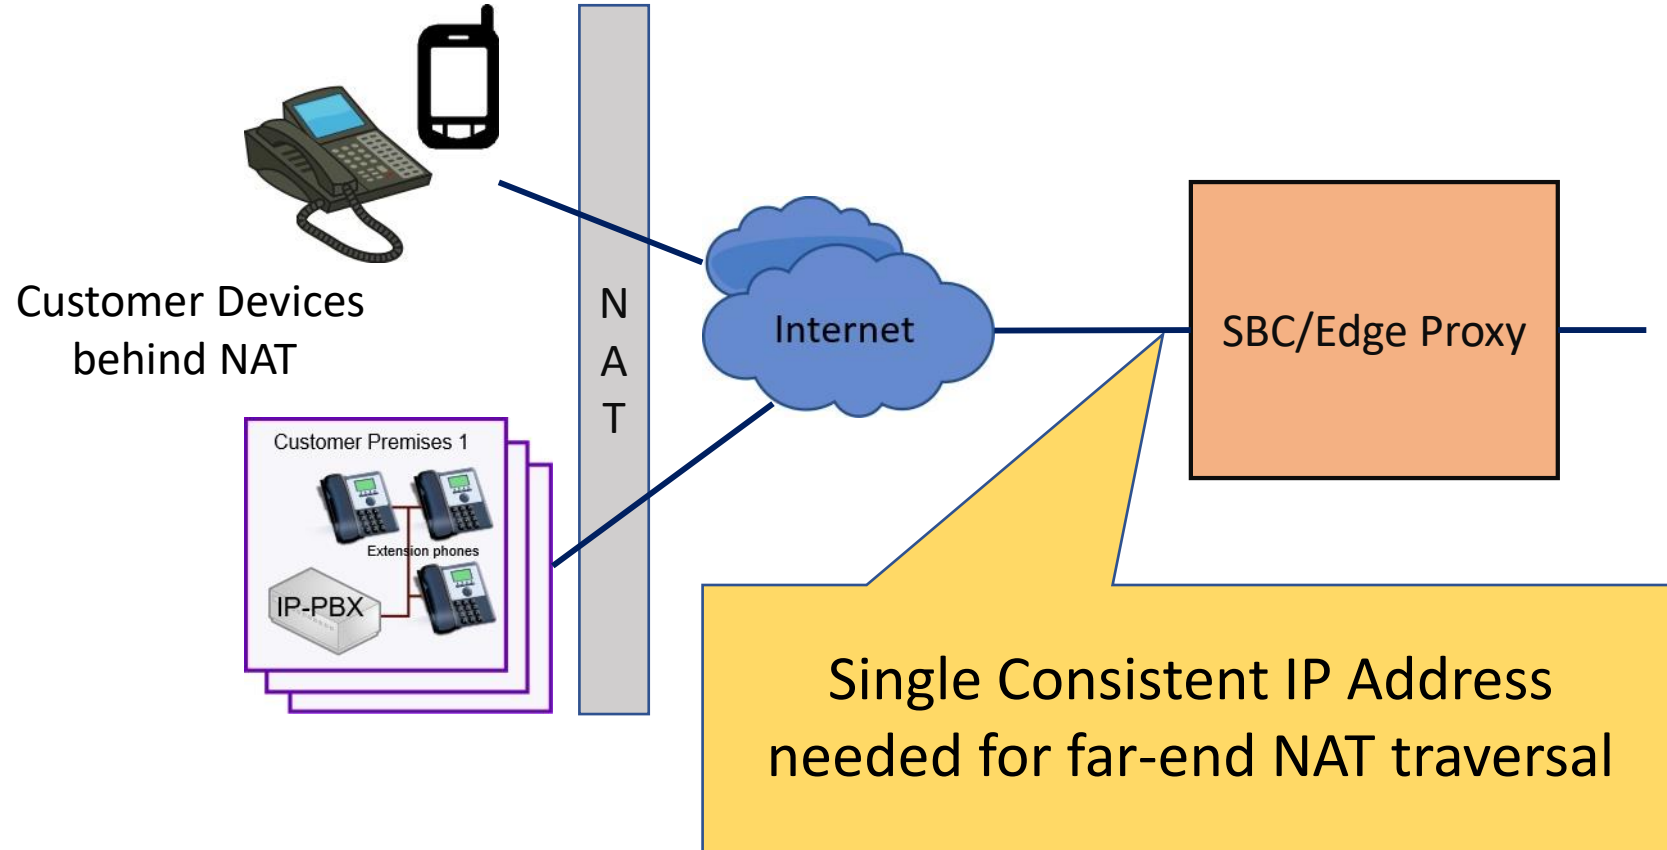
WebRTC

Kubernetes

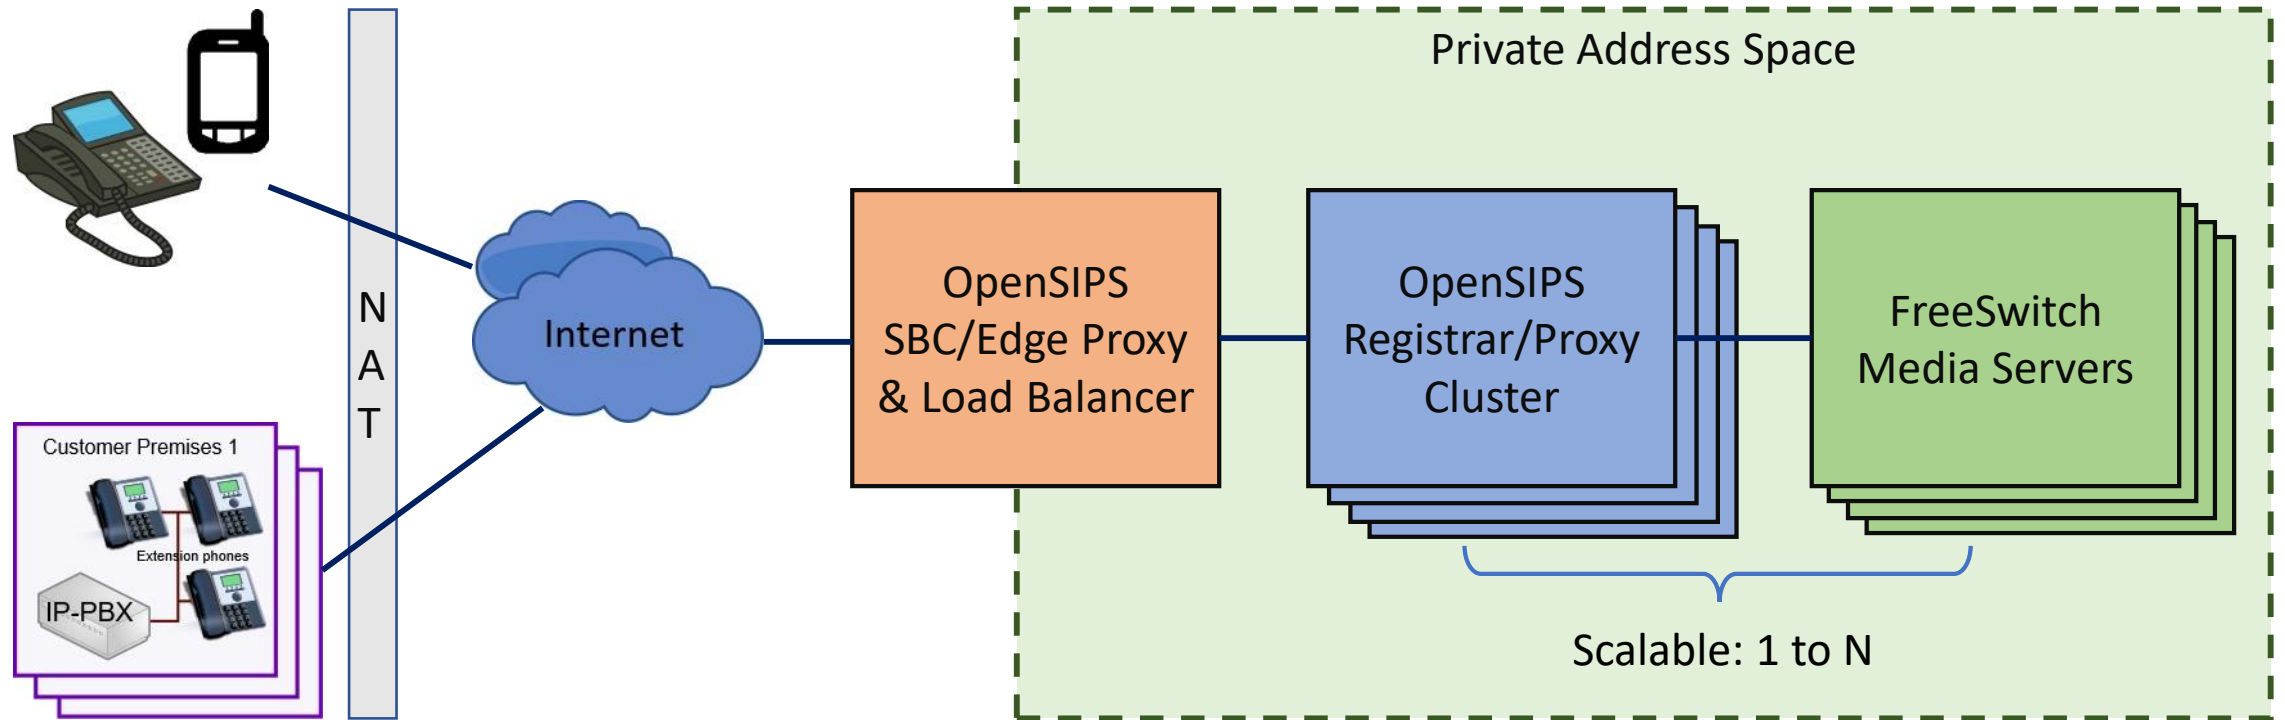
Clustering
Topology:
Initial Design
Inspiration



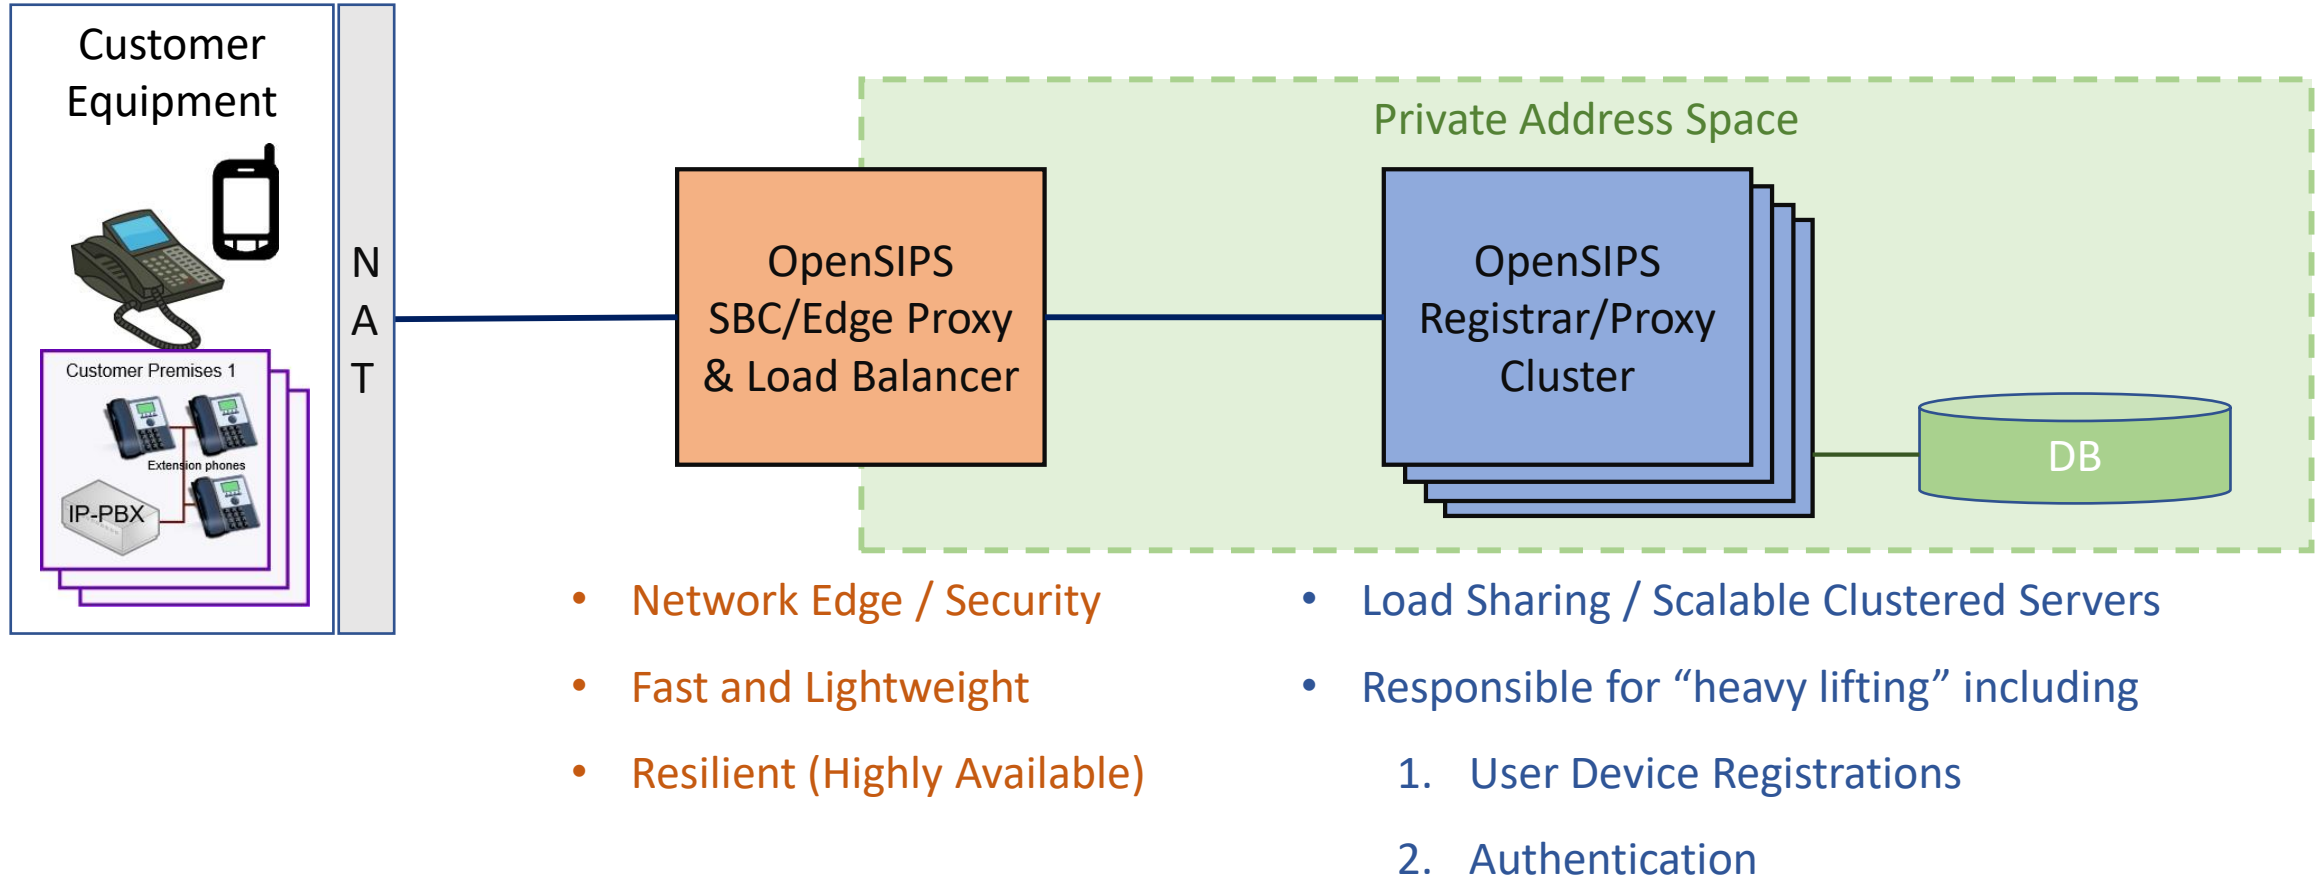
Why do we need the SBC / Edge Proxy?



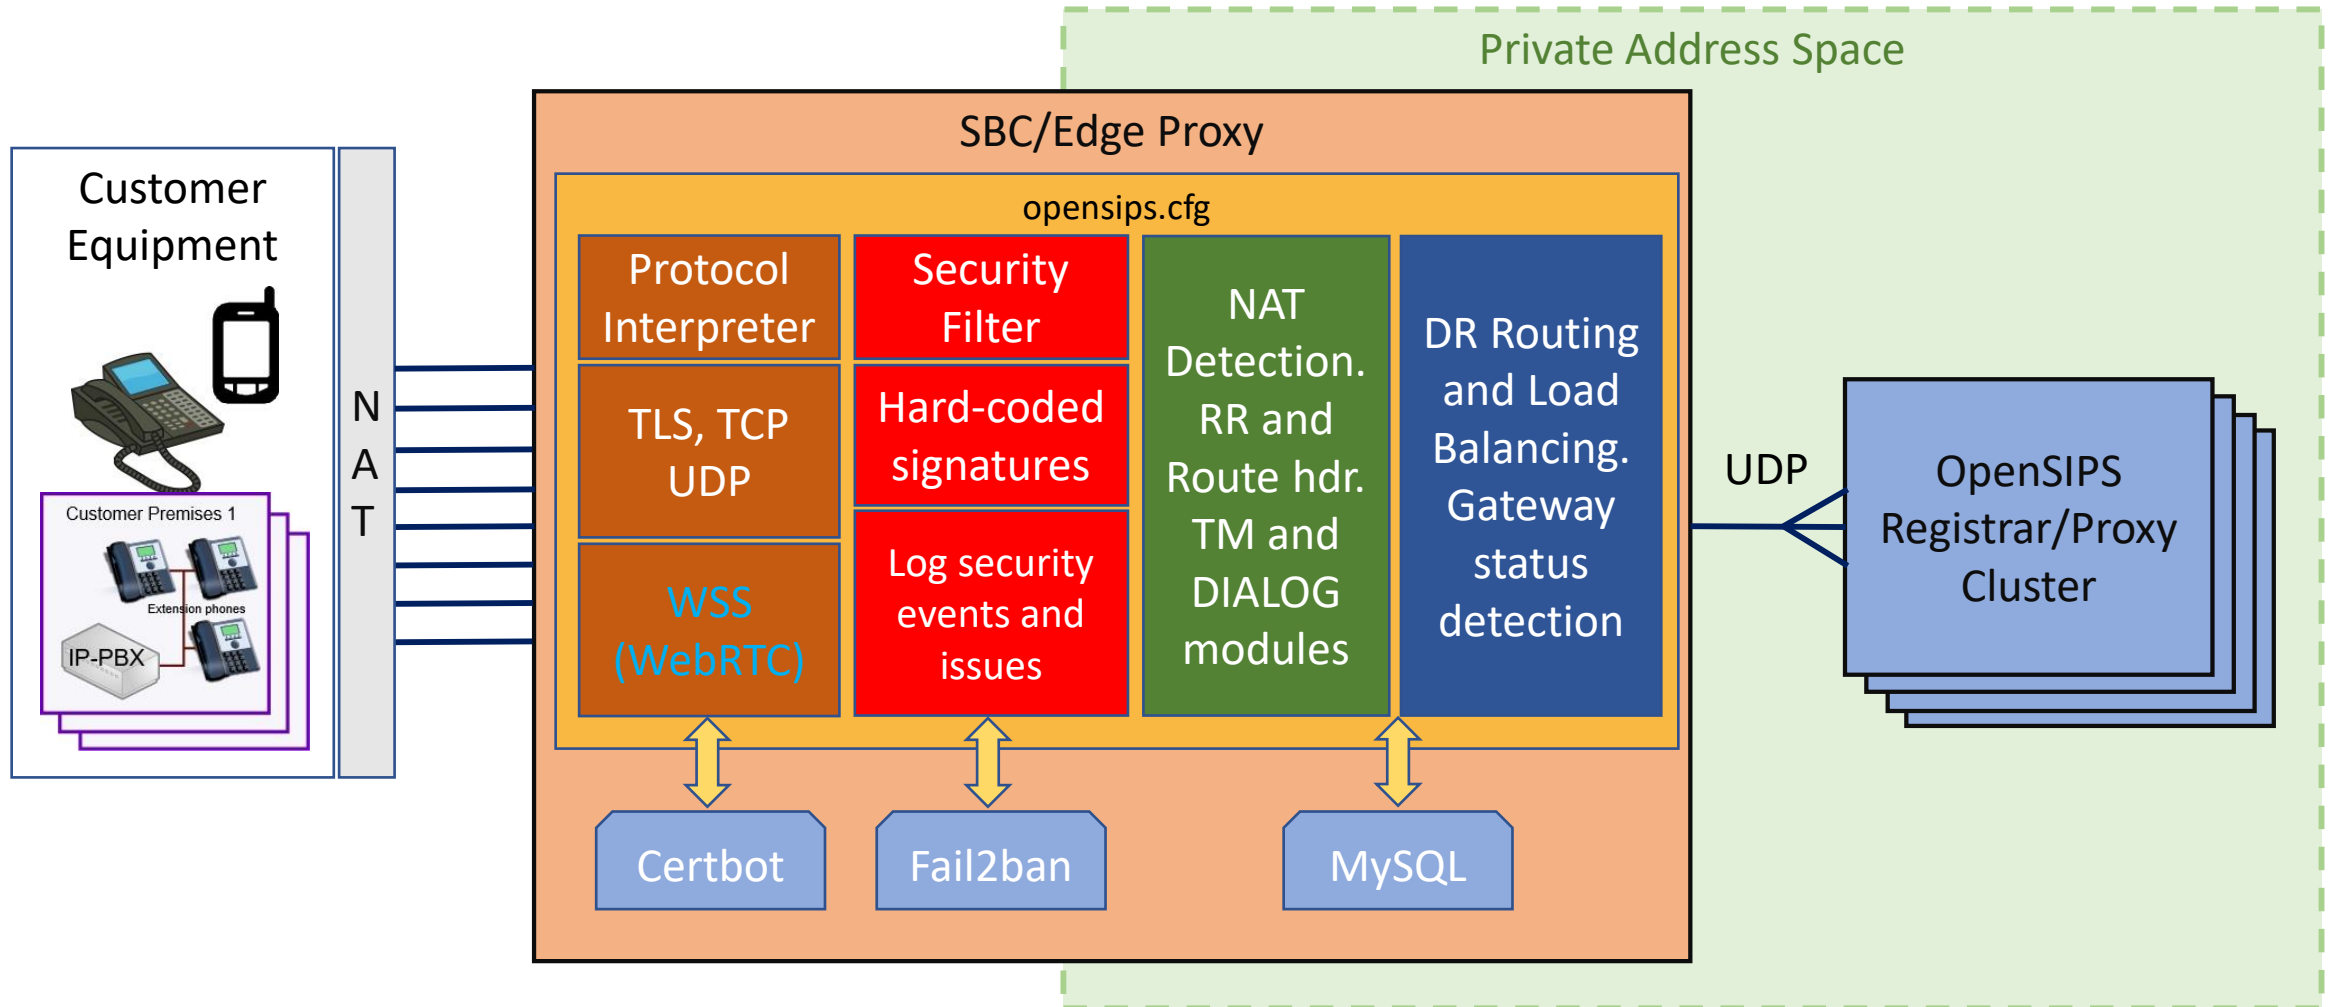
Clustered servers behind an Edge Proxy



OpenSIPS Server Roles



SBC / Edge Proxy - deconstructed



Risks when using SBC / Edge Proxy

- Single Point of Failure

HA Cluster using Pacemaker

BGP Network Connectivity

Future options – Dual DC, Anycast

- Traffic Bottleneck

Enable DNS Caching (module DNS_CACHE)

Optimise settings in opensips.cfg:

- open_files_limit
- children (worker processes)

Optimise and monitor pkmem & shmем

As Stateless as possible

As Stateless As Possible

- Problem:

Correct routing of sequential requests with minimal memory use

- Solution

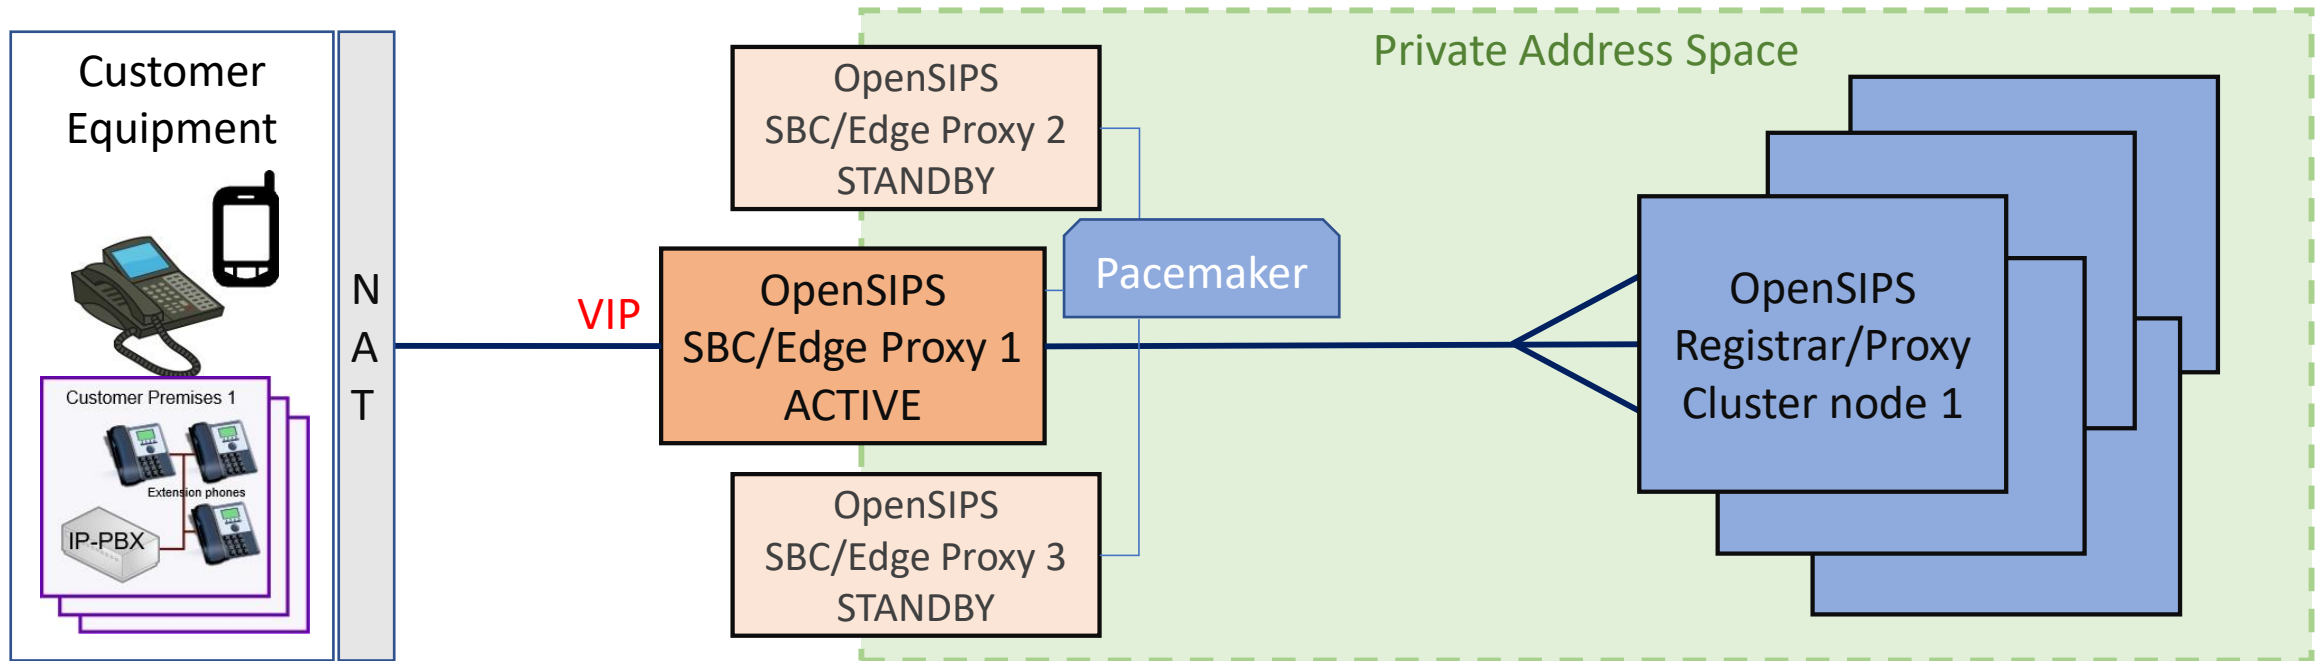
- a) Record-Route, Route and Contact headers
- b) Don't use Topology Hiding module

- Still to be resolved

TM and DIALOGUE modules

- a) Memory overhead?
- b) Are they even needed?

OpenSIPS: Application of Clusterer Module



- Shared DIALOG data
- Pacemaker controls VIP

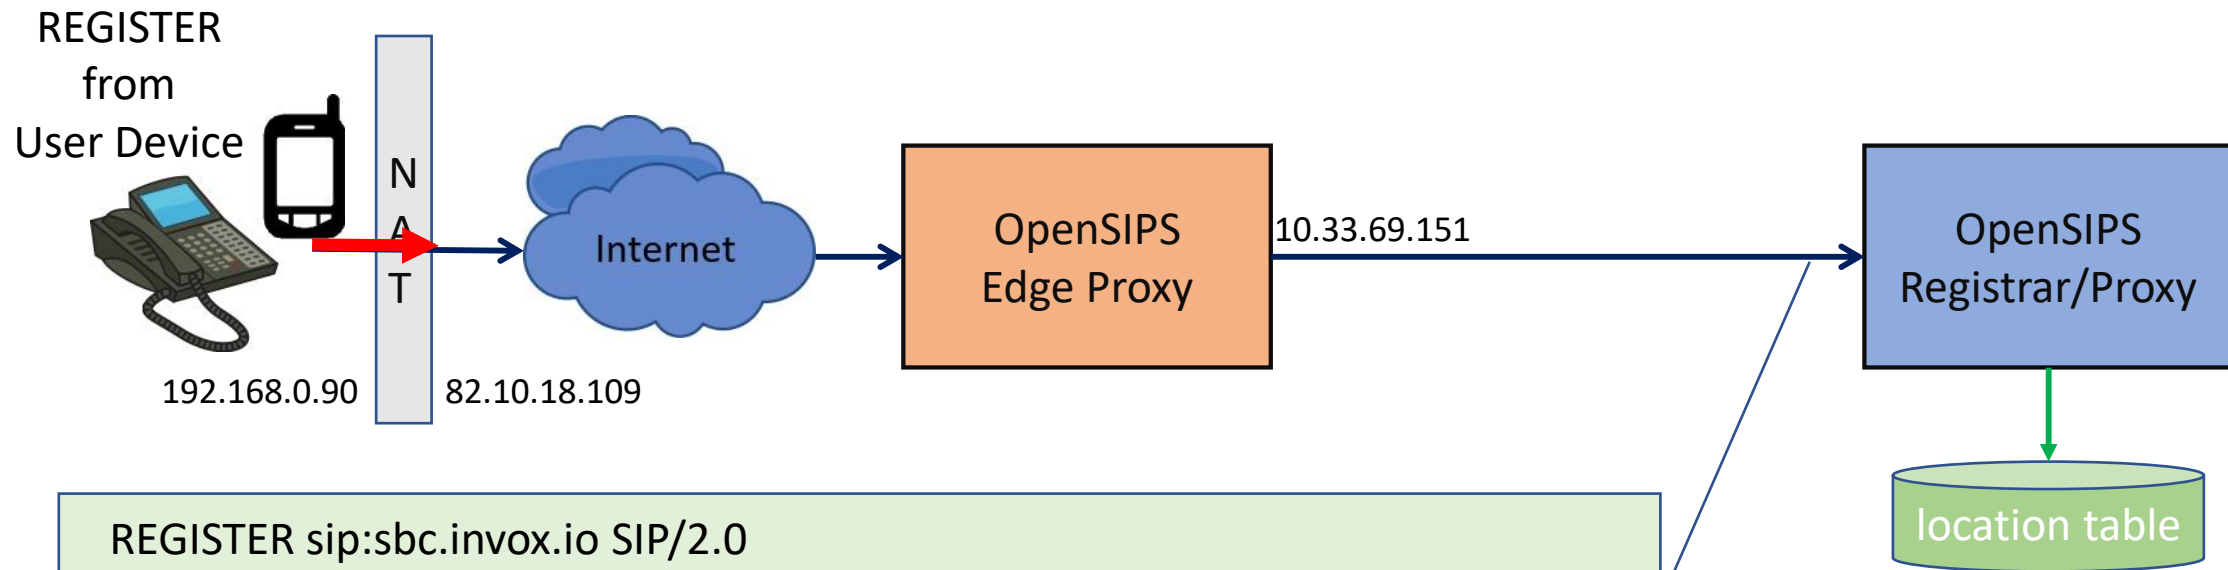
- USRLOC Full Sharing Cluster
- NAT Ping Sharing
- Scalable 1 to N Cluster Nodes



Topology- related Issues

Using Path headers for NAT'd user devices - 1

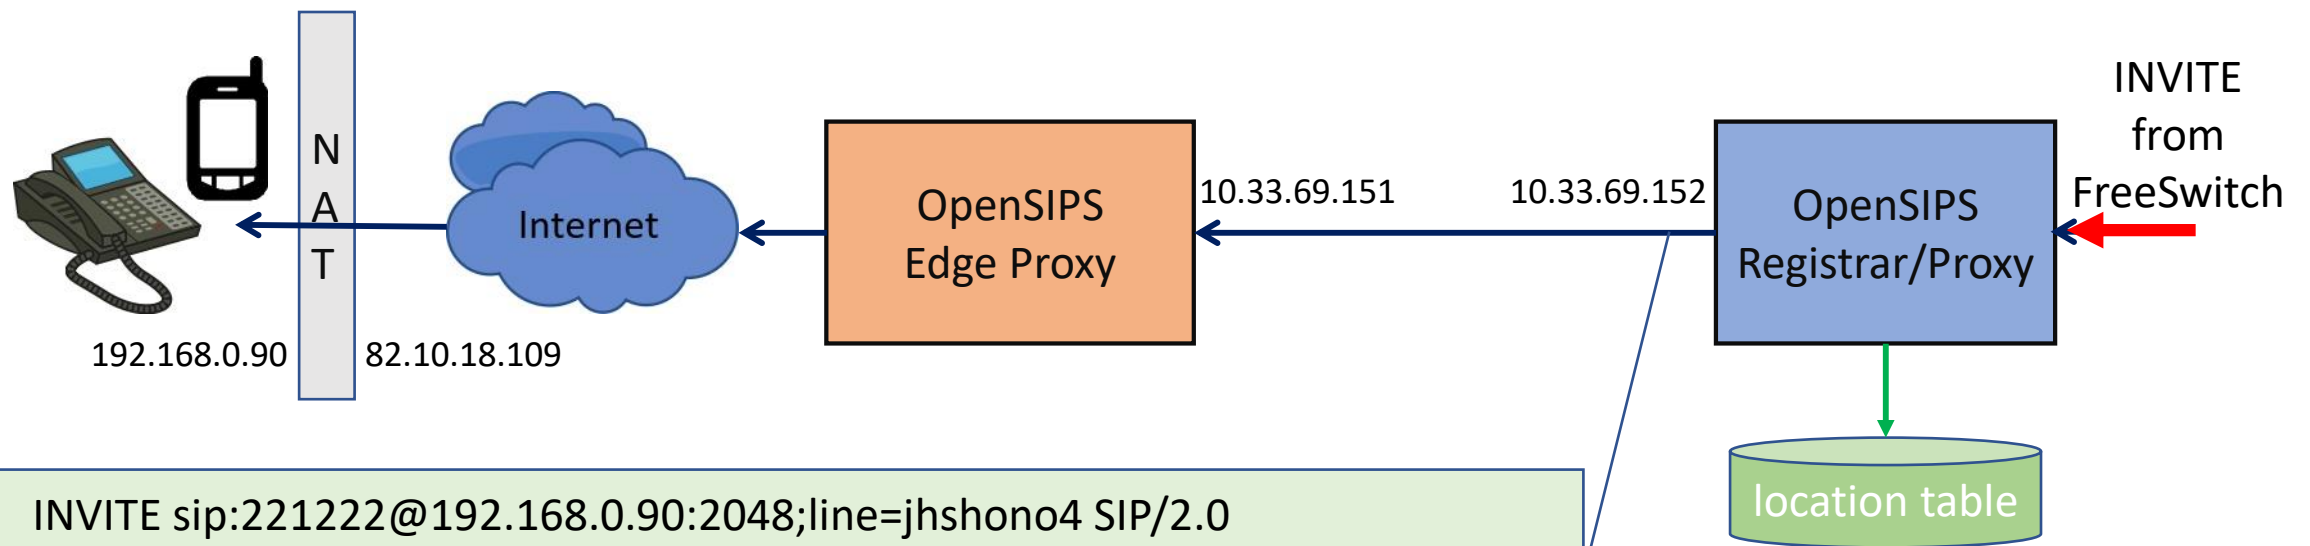
- Path Header is added to REGISTER request by our Edge Proxy
- Includes “received=” parameter when user device behind NAT



```
REGISTER sip:sbc.invox.io SIP/2.0
Via: SIP/2.0/UDP 10.33.69.151:5060;branch=z9hG4bK1351
Via: SIP/2.0/UDP 192.168.0.90:2048;received=82.10.18.109;branch=z9hG4bK.
Contact: <sip:221222@192.168.0.90:2048;line=jhshono4>
Path: <sip:mysbc@10.33.69.151;lr;received=sip:82.10.18.109:12202>
```

Using Path headers for NAT'd user devices - 2

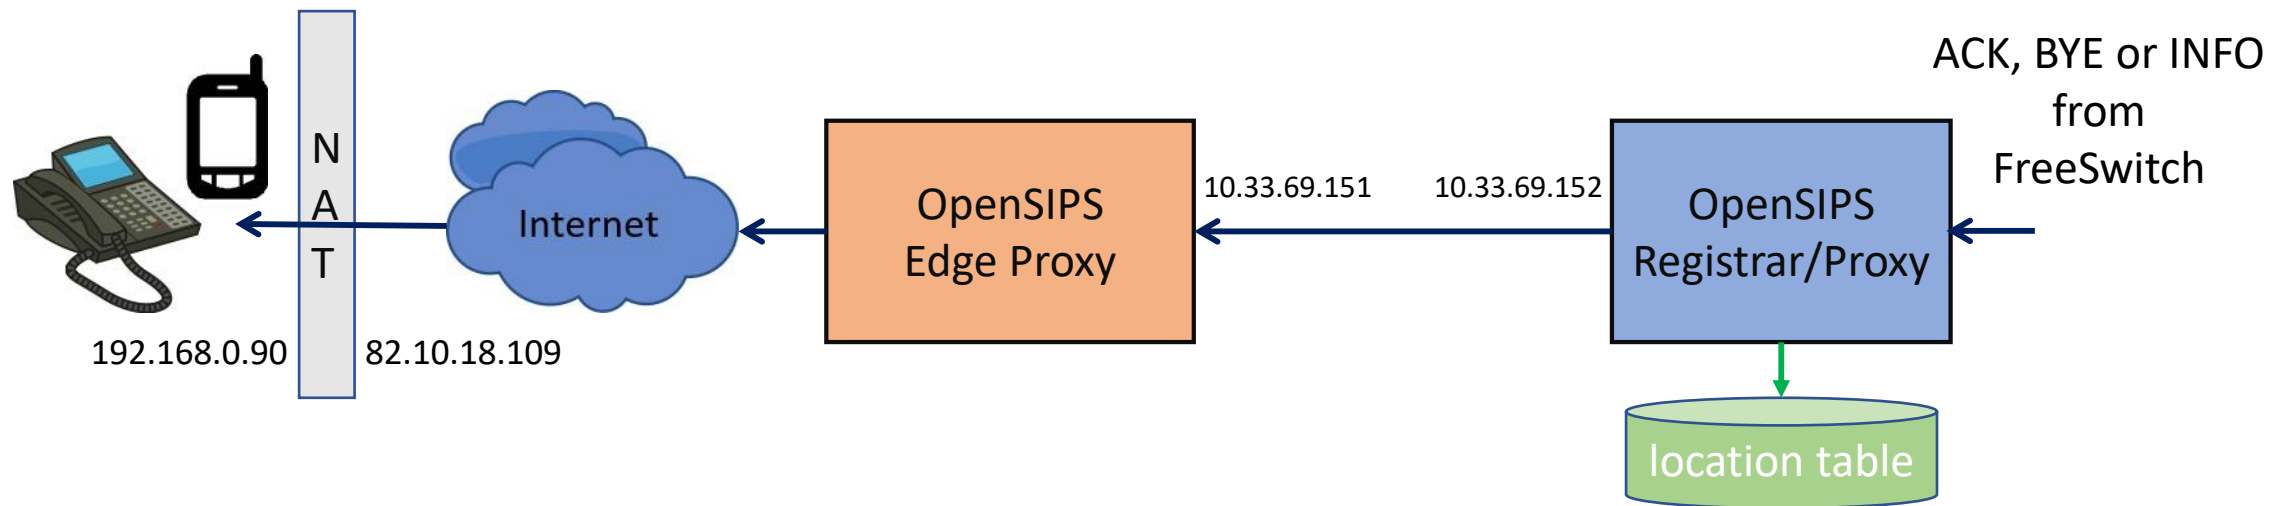
- In the Registrar Proxy, lookup() retrieves data from location table
- Registrar Proxy automatically adds a Route header based on Path



```
INVITE sip:221222@192.168.0.90:2048;line=jhshono4 SIP/2.0
Record-Route: <sip:10.33.69.152;lr;ftag=4eUm..;natua=10.33.69.152>
Via..
Route: <sip:mysbc@10.33.69.151;lr;received=sip:82.10.18.109:12202>
Other headers..
```

Using Path headers for NAT'd user devices - 3

- Sequential Requests such as ACK and BYE also need to call lookup



- If no lookup() in the Registrar/Proxy, there is no “received=” parameter
- No “received=” parameter means Edge Proxy cannot route Request

Using Path headers for NAT'd user devices

- Problem:

Edge Proxy must always receive a Route header with “received=” parameter when the request is going to a NAT'd customer device

- Solution

- a) During device registration, detect NAT and add a branch flag which gets stored in the relevant location table record
- b) During initial INVITE, do lookup(). If the NAT branch flag is present then add parameter “natua” to the Record-Route hdr.
- c) When handling sequential requests later, look for the “natua” parameter in the Route header. If found, do lookup() again

NAT'd user devices – code snippet

```
# Handle Loose_Routed packets
if (has_totag()) {
    if (loose_route()) {
        if (is_direction("downstream")) {
            if (check_route_param("natua=")) {
                remove_hf("Route");
                # Get best Path/Route info with location table lookup
                if (lookup("location"))
                    xlog("L_WARN", "    Lookup worked R-URI=$ru Src=$si\n");
                else
                    xlog("L_WARN", "    !!Loose-Routed $rm request failed lookup!! \n");
            }
        }
    }
}
```

Other Topology-related issues and snags

- Problem:

Load Balancer must send authenticated Requests to the same Registrar Proxy node that sent the Challenge

- Solution

- a) Use local CacheDB as short term, fast access memory store
- b) Load Balancing mechanism chooses node for initial request
- c) Push address of selected node (use source address as key)
- d) Request arrives with Digest Authentication
- e) Pop node address back from local CacheDB

Other Topology-related issues and snags


- Problem:

Customer devices dropping oversized INVITE requests;

MTU exceeded due to addition of Record-Route and Via headers

- Solution

- a) Strip out unnecessary extra headers added by FreeSwitch
- b) Could use Compact Headers
- c) Could use Topology Hiding, but only with a memory hit



docker

Introducing Docker

A Crash Course

The Benefits of using Docker

- A great Deployment Tool
 - Portability
 - Repeatability
 - Lightweight and Fast
- Encourages Good Practice
 - Microservices architecture
 - Well defined environment – O/S, Packages, Data, Configuration
 - Well defined build sequence
- A Basis for Scalability
 - PaaS – e.g. Kubernetes
 - IaaS – e.g. Amazon AWS

Some Docker basics

Image

- A file defining what will be in the container when it runs
- Official “base” images are available from Docker Hub

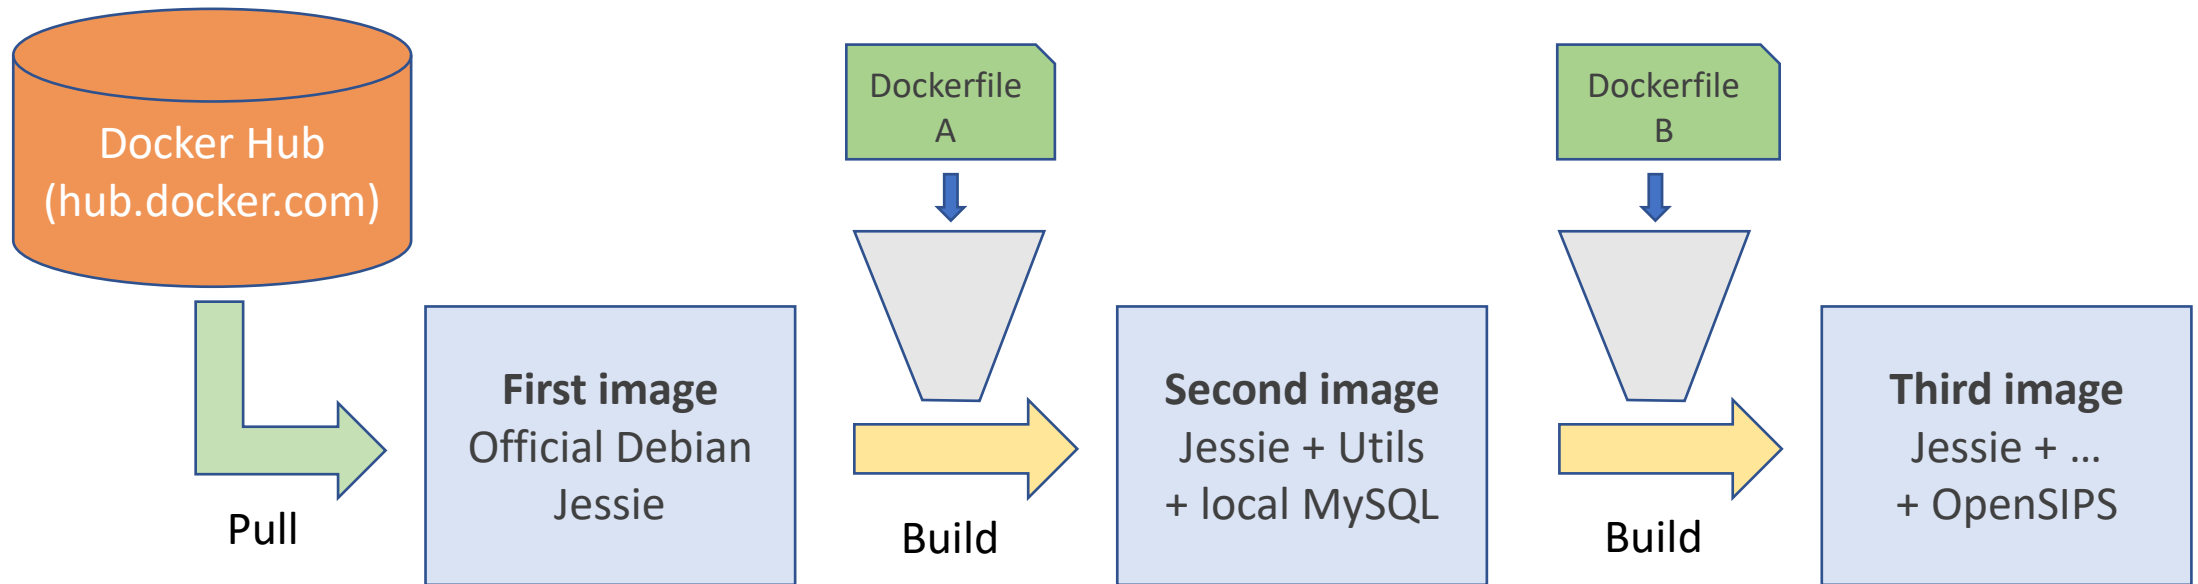
Docker Hub (hub.docker.com)

- A registry of publicly available Images – similar to github
- Pull images from Docker Hub into your own local store
- Push images to Docker Hub. Private storage areas are possible

Some Docker basics

Building your own Images

- New images are built, layer-by-layer, on top of existing one



Example of a Docker build file

```
FROM invozio/opensips-base-deb:v2.4.4
LABEL maintainer="john.quick@smartvox.co.uk"
WORKDIR /etc/opensips

# Copy the working cfg files for our OpenSIPS Server
COPY *.cfg /etc/opensips/

# Copy the service start bash script file to the image (overwrites any pre-existing version)
COPY service-start.sh /usr/local/sbin/
RUN chmod +x /usr/local/sbin/service-start.sh

EXPOSE 5060 5678

CMD ["/bin/bash", "-c", "/usr/local/sbin/service-start.sh"]
```

Some Docker basics

Image

- A file defining what will be in the container when it runs
- Like a virtual machine ISO image, but much lighter weight
- Defines run-time environment, application and data

Containers

- Are instantiated Docker Images
- Encapsulate the run-time environment and a deployed application

Docker basics – Running an image

docker run

..the command to run an image

-dti

..run in background with tty terminal

--env-file=runenvparms

..get parameter values from a file

--name registrar

..the container instance name

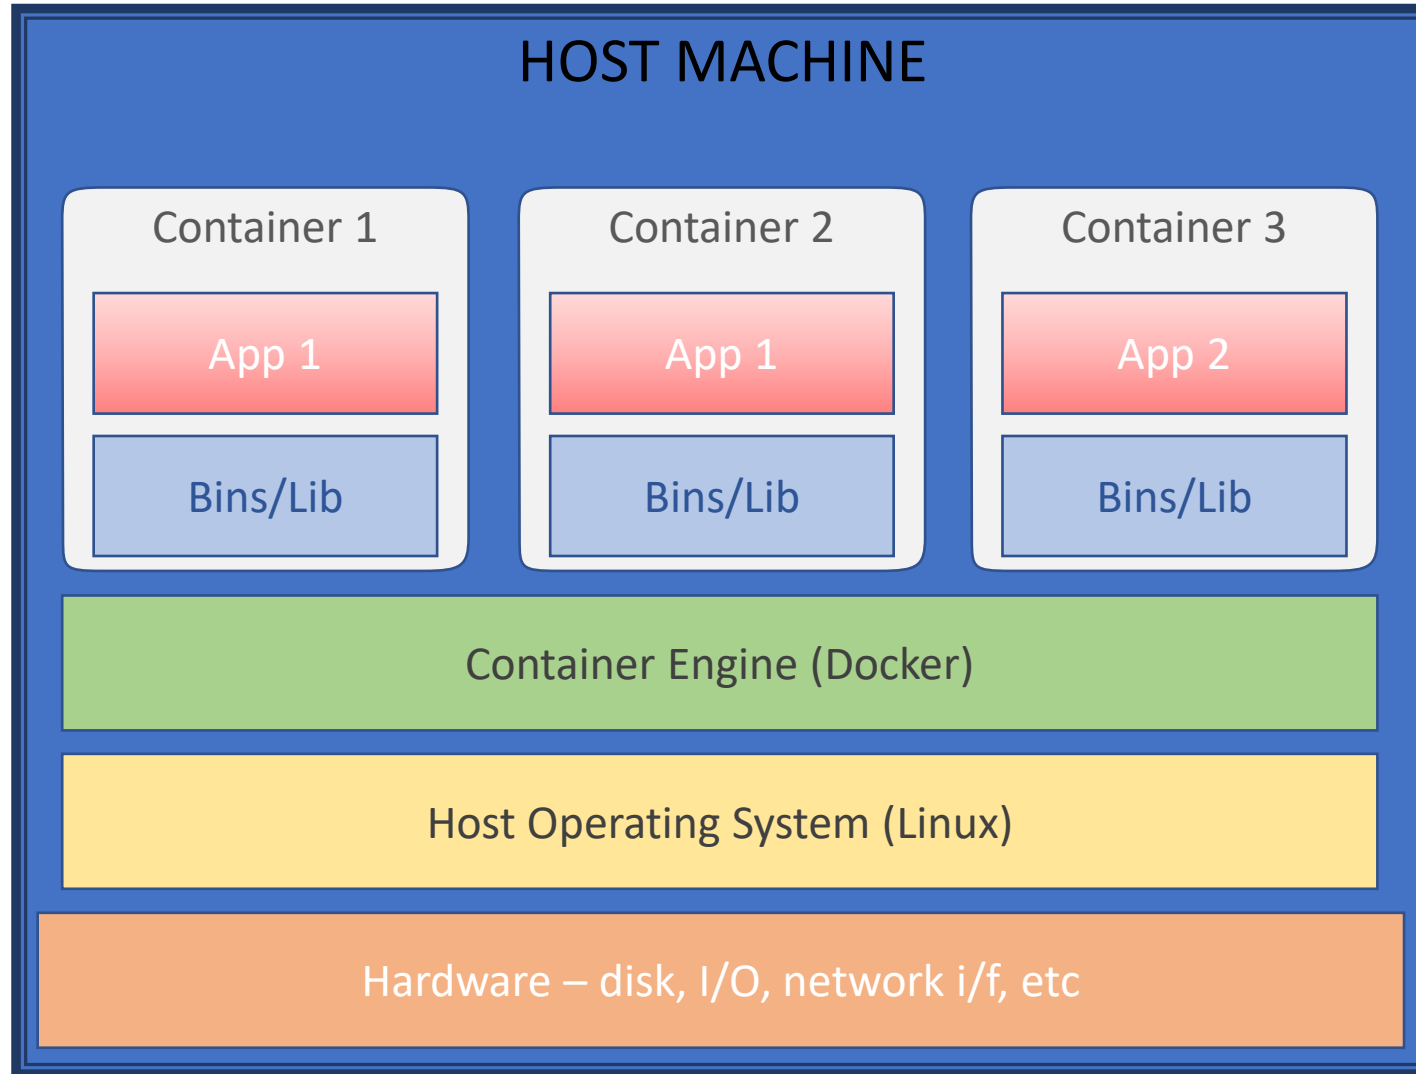
-p 5060:5060/udp

..expose ports


invoxio/opensips-reg


..use this image

Docker architecture



Docker networking (default bridging mode)



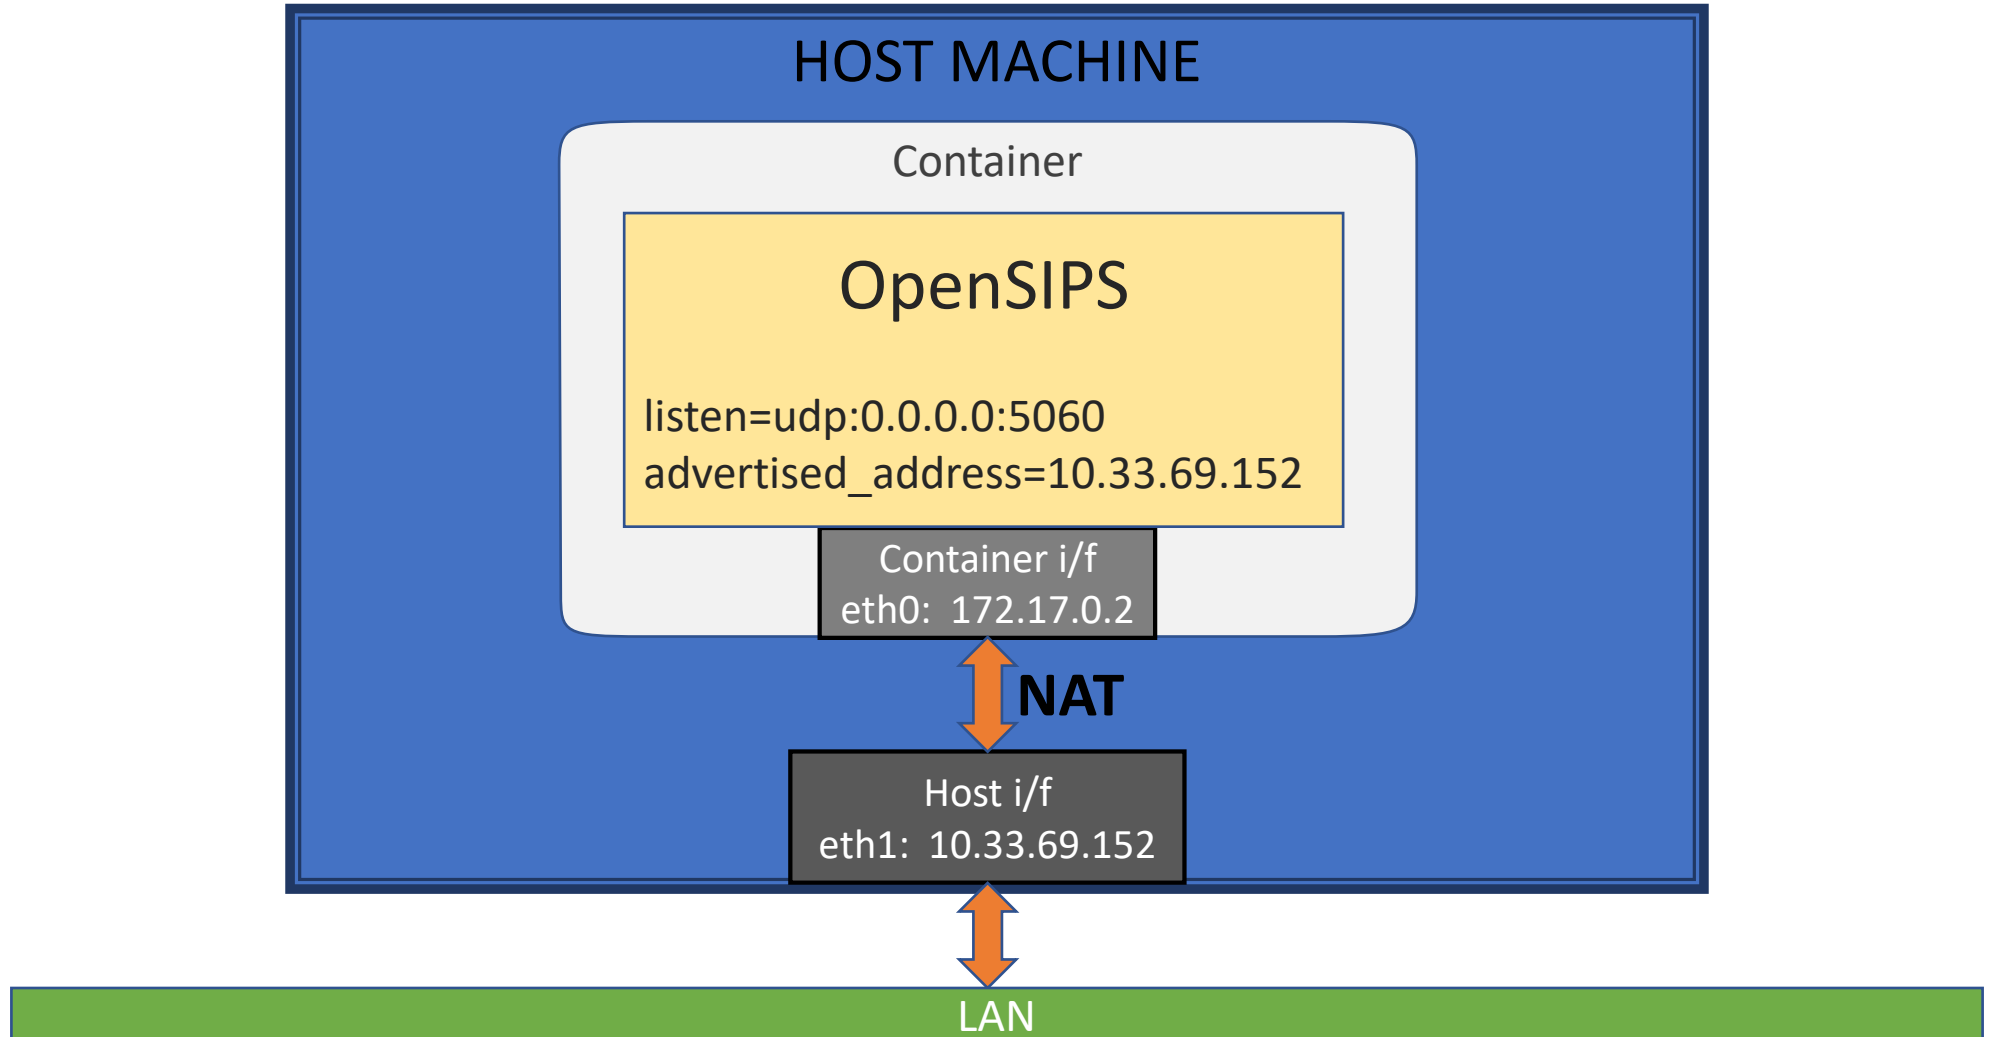


docker

Using Docker with OpenSIPS

Practical Application

Docker networking: OpenSIPS behind NAT



Docker Containers as Microservices

Idealised scenario

- One Application per Container
- Application runs in the foreground
- Application start is the last action during start-up
- The container will stop if the application stops

Docker Containers as Microservices

In Practice:

A Container is **primarily** used for one application

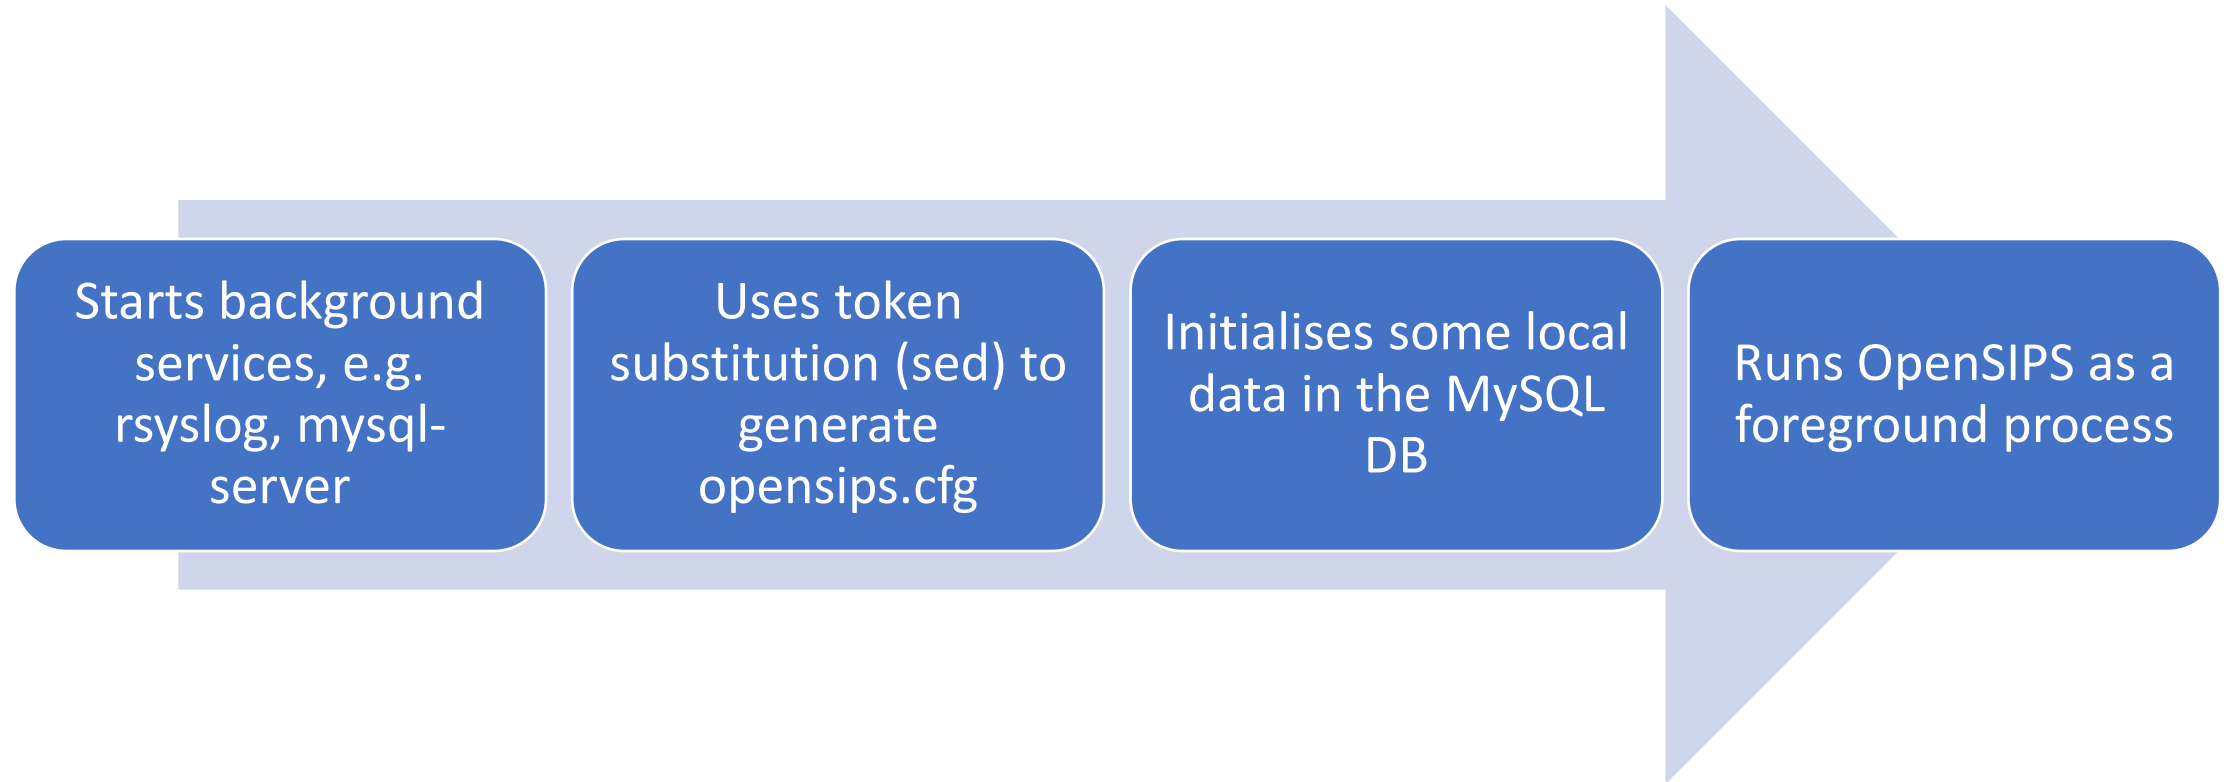
- e.g. I included a local instance of MySQL Server!

Last action during start-up runs a bash script

- The last action of the bash script runs the OpenSIPS app
- Is this Cheating??

Docker Containers as Microservices

In Practice: What my bash script does



Example start-up bash script file

```
#!/bin/bash
cat /etc/opensips/opensips.cfg.template | sed "s/%HOST_IP_ADDR%/${HOST_IP}/g" |
sed "s/%CLUSTER_NODE_ID%/${NODE_ID}/g" > /etc/opensips/opensips.cfg

service rsyslog start
service mysql start
sleep 2

if [ -e /etc/opensips/dbinit.sql ] ; then
    mysql --user=root --password=mypasswrđ opensips < /etc/opensips/dbinit.sql
    rm -f /etc/opensips/dbinit.sql
fi

/sbin/opensips -f /etc/opensips/opensips.cfg -F -m 640 -M 5
```



Docker- related Issues

Docker Containers as Scalable Cluster Nodes


Idealised Scenario

- Every container is identical
- Automatic neighbour discovery

OpenSIPS clusterer nodes

- Unique Node ID
- Special “seed” node
- Minimum of one static node address

Containerisation and Docker networking



- OpenSIPS is effectively behind NAT
- Docker re-organises iptables rules
- Fail2ban jails must use FORWARD chain
- Where to run Certbot & store TLS Certificates?
- Errors in opensips.cfg are hard to diagnose
- Problem diagnosis is generally more difficult
- Log files – redirect to host or syslog server



Running under Kubernetes


Work in progress

The role of Kubernetes



Define the physical resources: Nodes

Nodes are assigned roles: Masters and Workers



Define Apps: Containers running inside Pods

Pod name

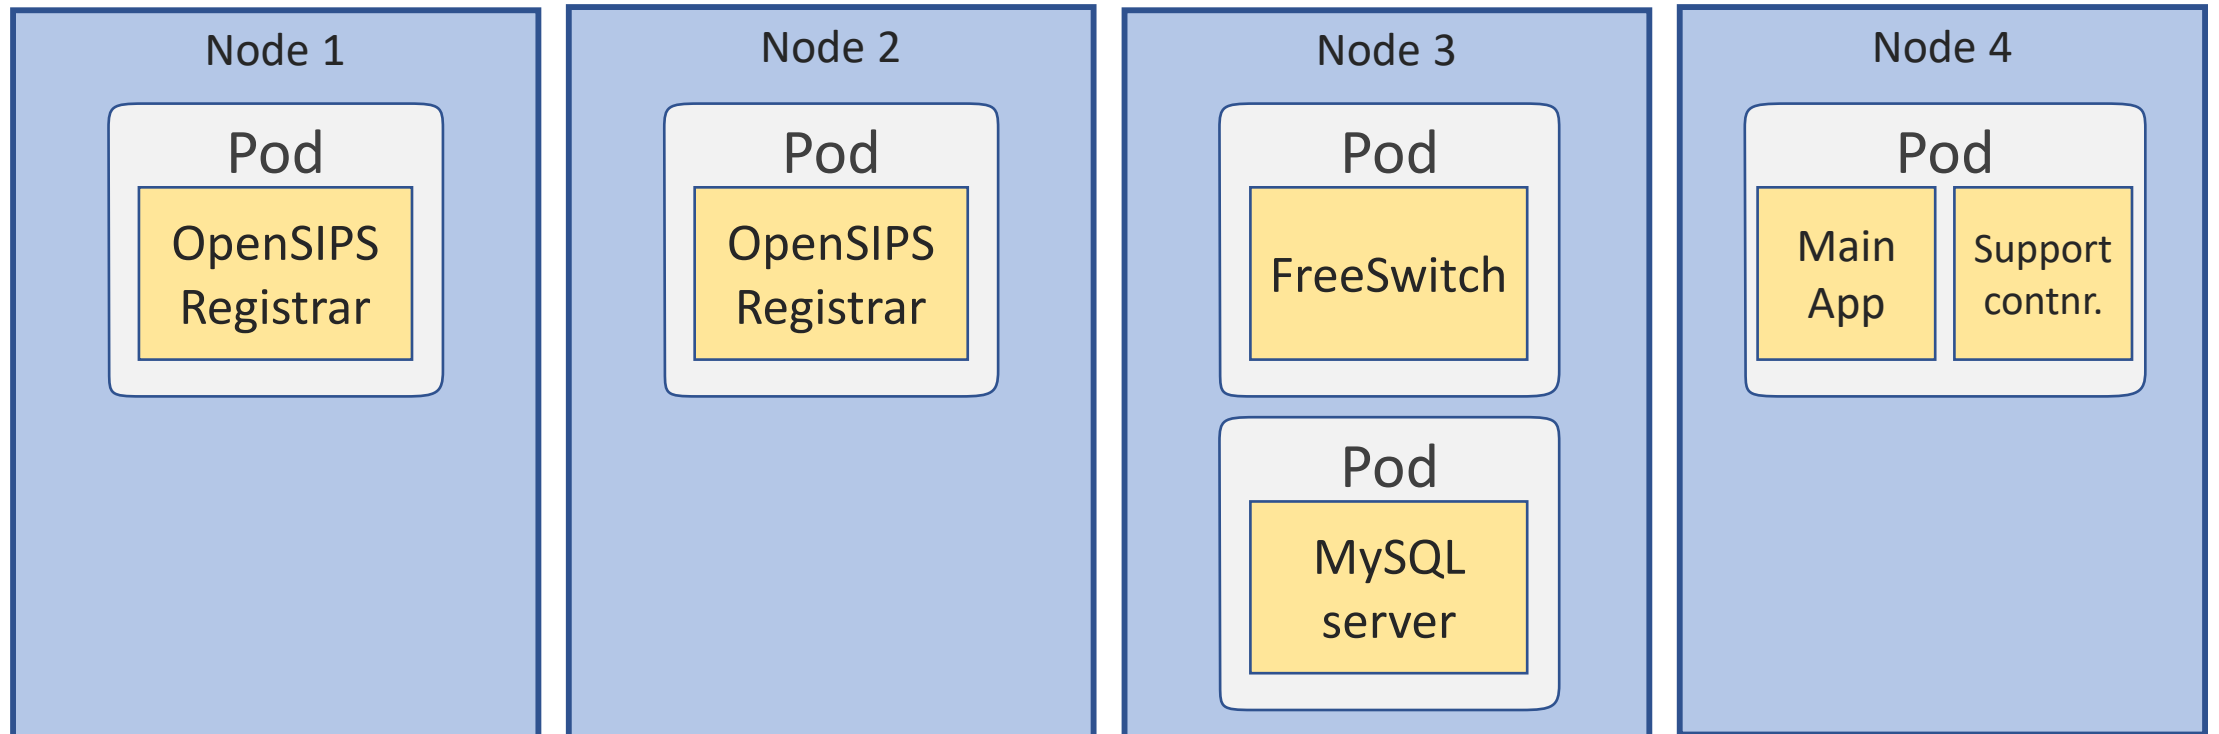
How many replicas

Container source image

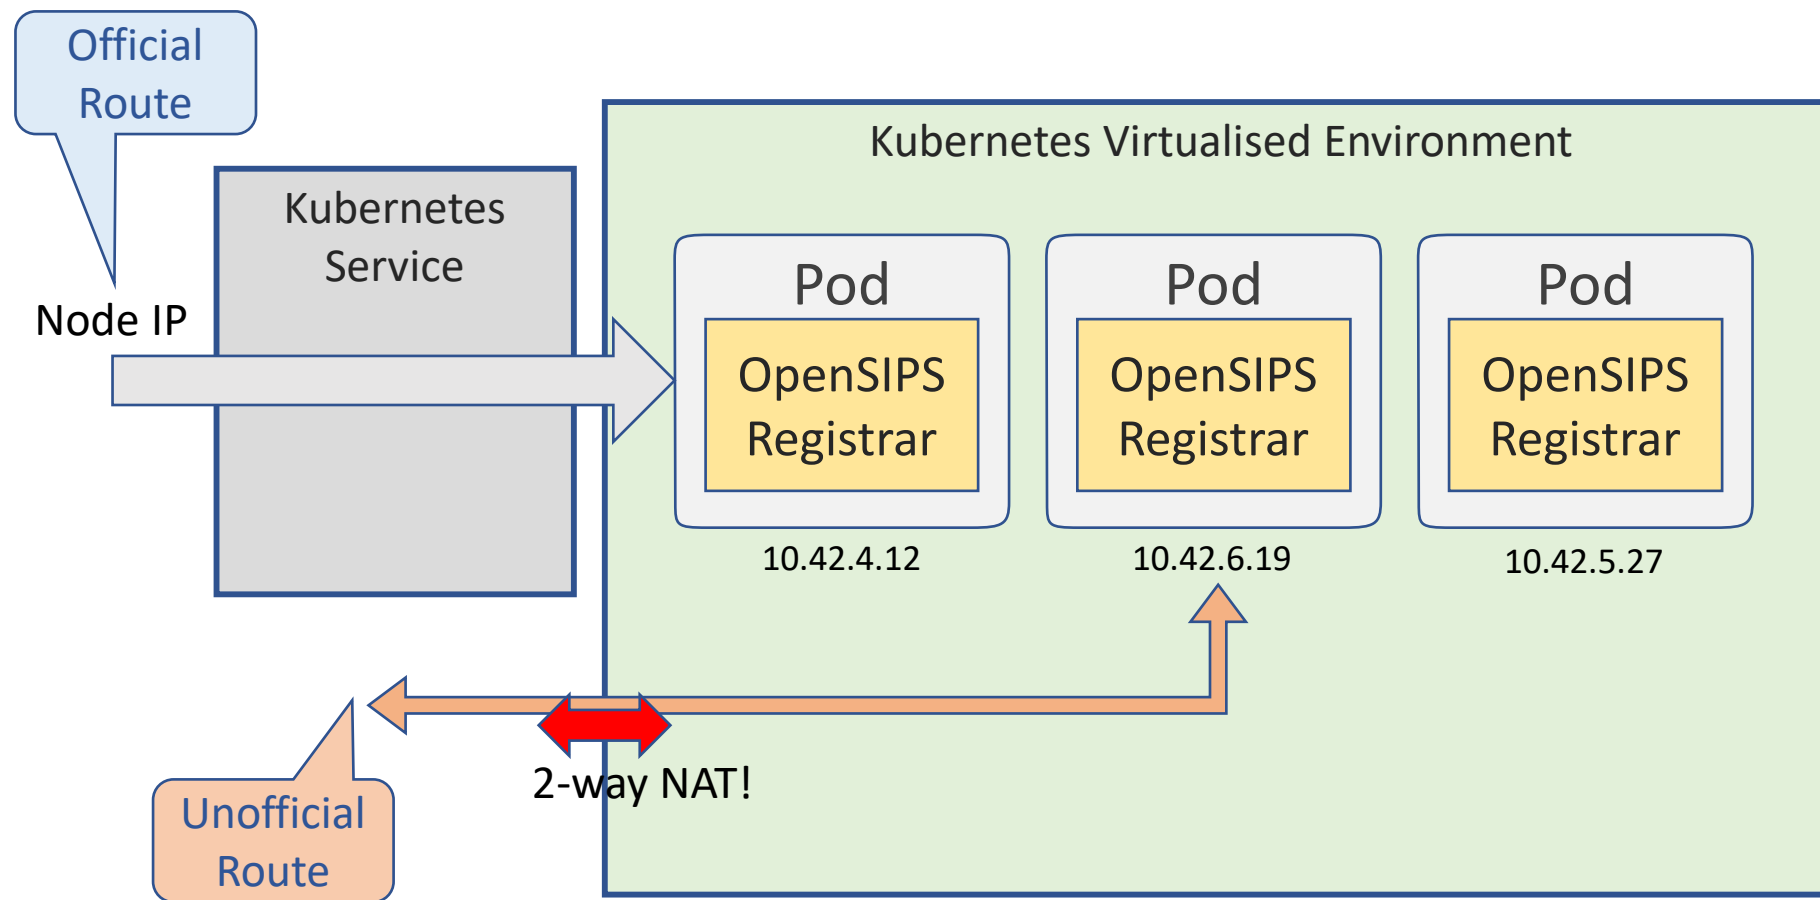
Ports to expose

Resource limits: Mem, CPU

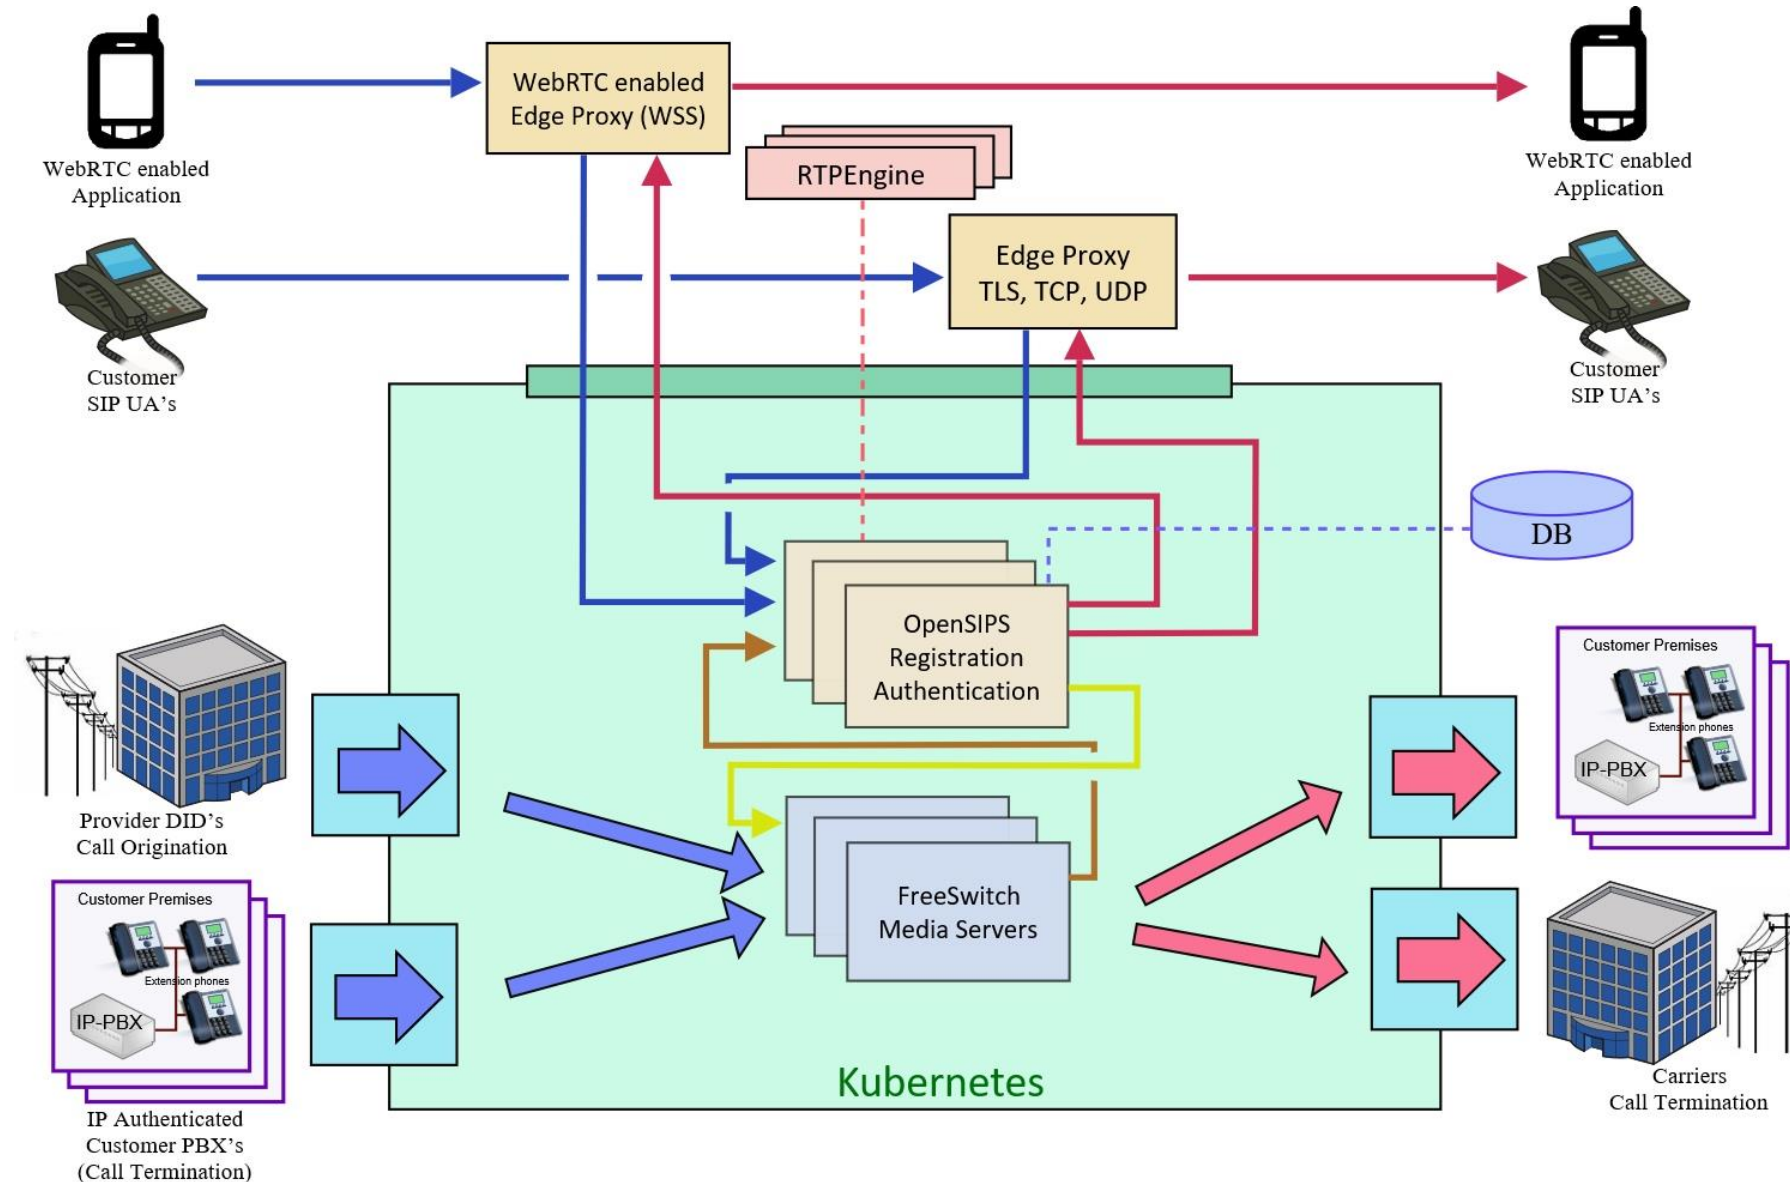
Metadata



Kubernetes Services and Networking



Where Things Stand Today



Thank You.
...and goodbye

John Quick

Smartvox Limited

Retiring this year 2019

Any Questions?

John Quick
Smartvox Limited