



Vlad Paiu OpenSIPS Summit 2025

Optimizing OpenSIPS Performance



- Why ?
- Practical Strategies forOptimizing OpenSIPS Performance
- Takeaways



Why optimize?

Why optimize ?



OpenSIPS offers the foundation for high performance

- o Written in C, fast:
 - SIP parsing
 - SIP processing & manipulation
 - core modules (TM, Dialog)
- Multi-process and asynchronous architecture
- Lean & Modular architecture

BUT...

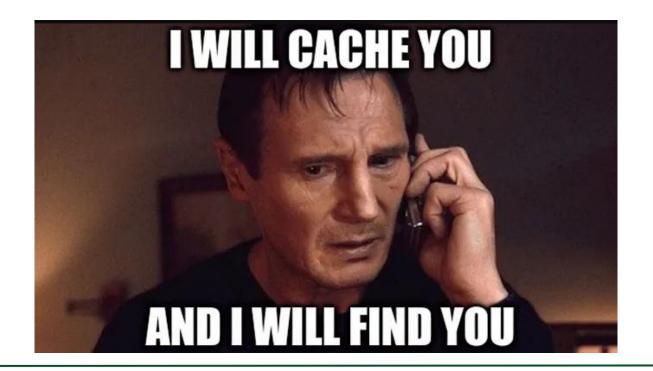
Why optimize?



- Suboptimal scripts can severely degrade OpenSIPS performance
- o Production environments are different than a controlled lab
- Unexpected traffic spikes can overwhelm untuned systems

Performance Optimisation Strategies







Many OpenSIPS modules use startup caching & runtime reload

- Domain
- Permissions
- Dispatcher
- Load Balancer
- Drouting



Wrong:

```
sql_query("SELECT client_id from my_clients where
ip_address='$si'","$avp(my_client_id)");
if ($rc > 0) {
          xlog("IP $si belongs to client $avp(my_client_id) \n");
}
```



Right:

```
loadmodule "permissions.so"

if (get_source_group( $avp(my_client_id) ) {
     xlog("IP $si belongs to client $avp(my_client_id) \n");
}
```



- Always question yourself if you are implementing things the right way
- Sometimes you might need to use DB Views to bring the data to the format that OpenSIPS knows to cache at startup

Cache all DB queries



 Sometimes you have custom tables that really don't fit into the predefined OpenSIPS tables

Do runtime caching for them

Cache all DB queries - Manual Approach



```
if (cache_fetch("local","client_id_$avp(client_id)",$avp(client_attr)) {
       xlog("Fetched attr $avp(client attr) for $avp(client id) from cache\n");
} else {
       sql query("SELECT client attr from my clients where id='$avp(client id'", "$avp(client attr)");
       if ($rc > 0) {
               # cache it forever
               cache_store("local","client_id_$avp(client_id)","$avp(client_attr)",0);
               xlog("Fetched attr $avp(client_attr) for $avp(client_id) from DB \n");
# For reloading : opensips-cli -x mi cache_remove local client_id_1
```

Cache all DB queries - SQL_Cacher Approach



```
modparam("sql_cacher", "cache_table",
"id=my_clients
db_url=mysql://root:opensips@localhost/opensips_3_6
cachedb_url=local://
table=my_clients
key=id
columns=client_attr
on_demand=1")
$avp(client_attr) = $sql_cached_value(my_clients:client_attr:$avp(client_id));
# For reloading : opensips-cli -x mi sql_cacher_reload my_clients 1
```

Try to Cache All External Dependencies



- HTTP queries
 - Manual Script caching approach
- DNS queries

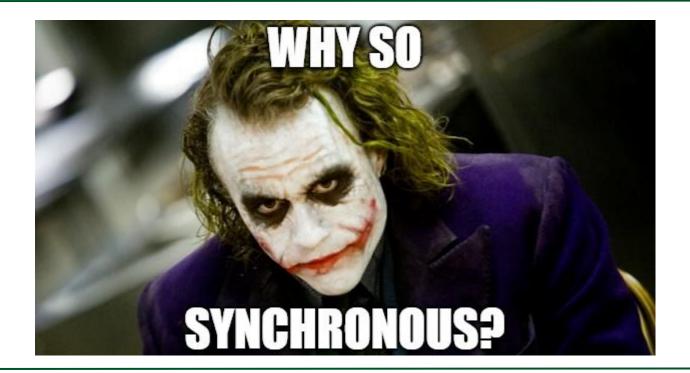
```
loadmodule "dns_cache.so"
modparam("dns_cache", "min_ttl",300) # new in 3.6
```

Try not to shoot yourself in the foot



- Maintain a balance between observability and performance
 - Avoid having a too verbose log
 - If you using Homer to replicate traffic, do it over UDP and do not use a DNS hostname for Homer itself
 - Monitor Homer aggressively and disable replicating if it's down
 - opensips-cli -x mi trace mode=off



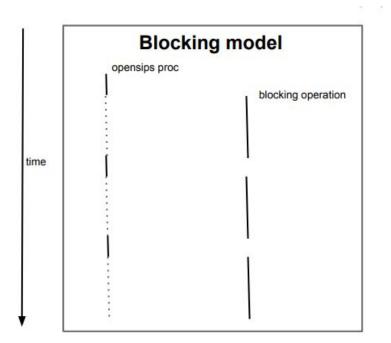


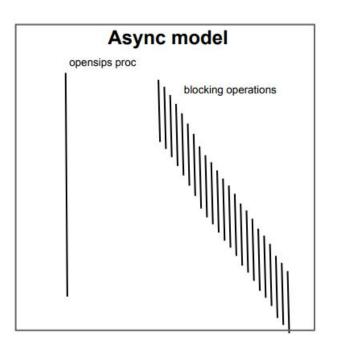


Aim for your OpenSIPS script to become CPU bound

- You never want to blockingly wait for an external dependency
 - o OpenSIPS processes will eventually get stuck waiting for the result
 - UDP/TCP queues will fill up
 - Packets & calls will start to drop









- dramatically improves scalability of OpenSIPS scripts
- better resilience to external service failures
- better usage of system resources

OpenSIPS waiting for RTPEngine







New in OpenSIPS 3.6

• Run async statements in onreply_route

• Async RTPEngine requests

Asynchronous RTPEngine Flow



```
route {
async(rtpengine_offer("$var(offer_flags)",$var(rtpengine_socket),$avp(sdp_body)),"RESUME_RTPENGINE_OFFER",2);
route[RESUME_RTPENGINE_OFFER] {
       if ($avp(sdp_body) != NULL) {
                # apply the changes in the SDP
                route(relay);
       } else {
                # failed to engage RTPEngine
                t_reply(503,"Unavailable");
                exit;
```

Asynchronous RTPEngine Flow



```
onreply_route[reply_processing] {
async(rtpengine_answer("$var(answer_flags)",,$avp(r_sdp_body)),"RESUME_RTPENGINE_ANSWER",2);
route[RESUME_RTPENGINE_ANSWER] {
       if ($avp(r_sdp_body) != NULL) {
                # apply the changes in the SDP
                # default action in reply context is for it to be relayed
       } else {
                # failed to engage rtpengine
                t cancel branch();
```

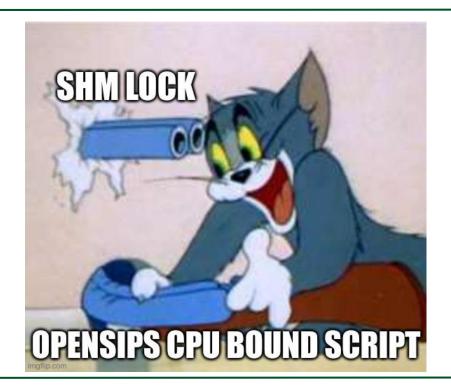
Asynchronous RTPEngine Flow



```
if (is_method("BYE")) {
    # launch the rtpengine_delete() and don't care whether it succeeds for not
    launch(rtpengine_delete());
```

My script is now CPU bound!



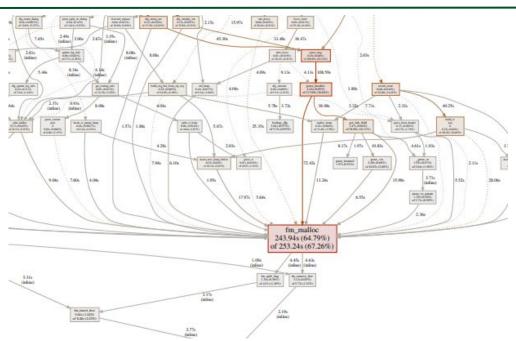




The shared memory allocator uses a global lock

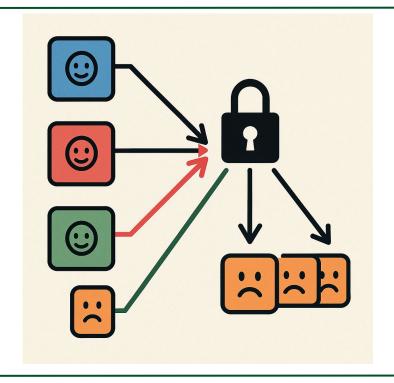
 In a high CPS & high CPUs environment, in combination with a very complex OpenSIPS script, this will lead to contention & starvation





https://www.opensips.org/About/PerformanceTests-3-4





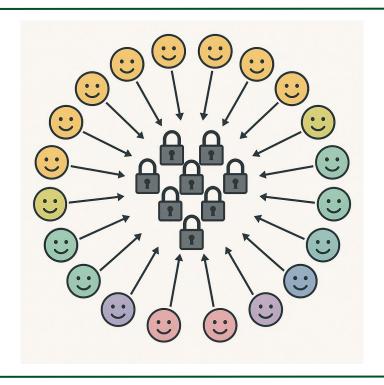


- Try to avoid contention
 - Use proper variable scope
 - \$var is per process allocated memory
 - The rest (\$avp, \$dlg_var, \$shv) are allocated in shared memory
 - When DB querying, if you know you'll get a single row:
 - Use sql_select_one(), which allows you to store output in a \$var



- OpenSIPS 3.6 adds a new F_PARALLEL_MALLOC allocator
 - Splits the total shared memory size into 32 separate memory chunks
 - Each memory chunk has its own allocator lock
 - When a process allocates shared memory it picks a random memory pool
 - When a process frees up memory, it puts it back in the original memory pool
 - o Can be set to be used at runtime via:
 - opensips -f etc/opensips.cfg -m 4096 -a F_PARALLEL_MALLOC -k F_MALLOC



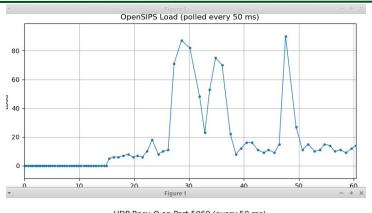


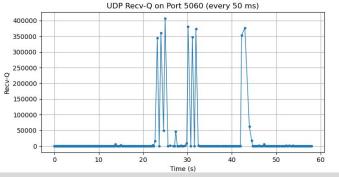


- Testing with a TM/DIALOG/ACC memory intensive script:
 - Setting 50 different \$AVPs at INVITE time
 - Setting 20 different \$dlg_vals at INVITE time
 - Pushing 50 values to ACC extra fields
- i7-1355U 6 cores / 12 threads
- No logging, SIPP with call duration of 10 seconds, running on the same box as OpenSIPS
- Plotting the load:load and UDP buffer queue size

F_MALLOC - 1800 CPS

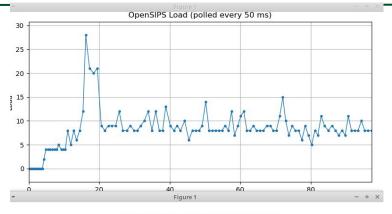


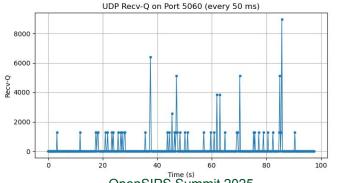




F_PARALLEL_MALLOC - 1800 CPS

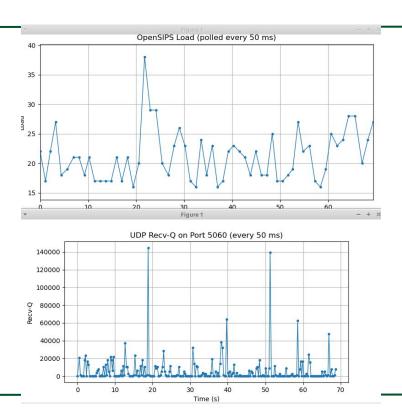






F_PARALLEL_MALLOC - 3000 CPS







 F_PARALLEL_MALLOC allows you to squeeze ~ 2x more CPS compared to F_MALLOC in high contention/starvation environments

 Much more resilient to temporary spikes and retransmission floods

What else can we improve?

Memory footprint & Reload Times



- What about the memory footprint ?
- What about the reload times
 - O How fast can you actually apply changes?

Caching 10 Million 10D prefixes



Use drouting module for caching the prefixes

- SHM Memory Usage:
 - o 10 GB; peak memory usage 20GB when doing a dr_reload
- dr_reload time on a local SQL:
 - o 140 seconds

Caching 10 Million 10D prefixes



- OpenSIPS 3.6 adds a new trie module
 - o https://opensips.org/html/docs/modules/3.6.x/trie.html
- SHM Memory Usage :
 - 4 GB = ~ 2.5x less memory usage vs drouting prefix caching; 5x less memory usage since there is no need for a full reload to apply changes
- Deleting numbers is instantaneous :
 - opensips-cli -x mi trie_number_delete partition_name=dids number=["1234567890","0987654321"]
- Adding numbers is instantaneous:
 - o opensips-cli -x mi trie_number_upsert partition_name=dids number=["1234","5678"]
 attrs=["1234attrs","5678attrs"]

Takeaways

Takeaways



Performance is in the hands of the OpenSIPS user/script writer

 OpenSIPS 3.6 has many improvements which allow you to squeeze as much performance as possible



Questions?

- Vlad Paiu
 - OpenSIPS Project: www.opensips.org
 - o Email: vladpaiu@opensips.org